

NANDISHA D

Generative AI Engineer | LLM Systems, RAG Pipelines & Agentic AI

Bangalore, India

+91-8088938270 | naik.nandishd@gmail.com | [linkedin.com/in/nandishd](https://www.linkedin.com/in/nandishd) | github.com/NandishNaik01 | nandishnaik.netlify.app

Generative AI Engineer with 2+ years shipping production **LLM systems, Retrieval-Augmented Generation (RAG) pipelines, multi-agent workflows, and Model Context Protocol (MCP) tooling**. At **Harvey.ai**, architected an end-to-end AI conversation pipeline for document-grounded legal research with streaming inference, hybrid retrieval, and hallucination-mitigation guardrails. Authored **6 Claude Code AI automation skills** cutting engineering cycles from weeks to days. Founded **Forge41**; built **GenAlpha CLI** – open-source tool converting any API repo into an **MCP server** via Temporal workflows and Python AST parsing. Rated **Above Expectations** (highest rating) across all performance reviews.

EXPERIENCE

Think41

Jun 2024 – Present

Full-time Employer – Career progression: Intern → SDE-I → SDE-II

Bangalore, On-site

- **Intern (Jun 2024 – Jan 2025)**: Ramped up on **Python, FastAPI, Django, React** full-stack delivery. Architected **RQ** (LangChain-based conversational PRD generator with LiveKit real-time collaboration) and **MUCAI** (production multi-agent conversational AI for HR/IT/PM with per-domain memory isolation, deployed on Docker + NGINX).
- **SDE-I (Feb 2025 – Mar 2026), client Atomicwork**: Authored **6 Claude Code AI automation skills** – LLM-executable workflows encoding the full integration lifecycle; cut per-integration delivery **from weeks to under one week**. Adopted team-wide and by the PM team for autonomous AI-agent POC integrations via **Devin** and Claude Code; **accepted into Atomicwork's official GitHub org**.
- Led sprint planning/delivery for the Integration Pod (**Agile/Scrum**): shipped **16 work items across 47 merged PRs in 6 months, zero rollbacks**; rated **Above Expectations**. Treated prompts as production code – versioning, structured evals, regression testing, telemetry.
- **SDE-II (Apr 2026 – Present), client Harvey.ai**: Built end-to-end **RAG pipeline** for document-grounded legal research: semantic chunking, **OpenAI text-embedding-3-large, pgvector** hybrid BM25 + vector retrieval, re-ranking, context assembly, and low-latency **streaming LLM inference** at production scale. Integrated **Ansarada DMS** via **Temporal** durable workflows (OAuth 2.0 lifecycle, Redis caching, dead-session recovery).
- Sole architect of Harvey's **multi-agent workflow automation engine** – led HLD, LLD, and first prototype **within one week**; orchestrates branching logic, tool use / function calling, and state persistence on **Temporal**. Engineered CoT prompting, structured output, **LLM-as-a-judge** evals, PII redaction, and prompt-injection defense; applied **Anthropic prompt caching** to cut per-request inference cost.

PROJECTS (OPEN-SOURCE AI INFRASTRUCTURE)

Harvey.ai POC – Voice-Native Legal RAG Assistant | *LiveKit, Pipecat, Google ADK, Deepgram, ElevenLabs, Turbopuffer* 2026

- Voice-native RAG proof-of-concept for legal research: real-time audio pipeline on **LiveKit + Pipecat** with **Google ADK** multi-agent orchestration, **Deepgram STT / ElevenLabs TTS** for low-latency conversational turns. File-type-aware chunking (Reducto OCR for scanned/image documents), **OpenAI text-embedding-3-large** embeddings in **Turbopuffer**, hybrid BM25 keyword + semantic retrieval with citations surfaced in the Word add-in.

GenAlpha CLI – API-to-MCP Server Generator | *Python, FastAPI, Next.js, Temporal, PostgreSQL, FastMCP* 2025 – Present

- Open-source **MCP server + CLI generator**: converts any API repo into a production **Model Context Protocol server** via OpenAPI detection + Python AST route extraction – exposing arbitrary APIs as AI-agent tools for Claude, ChatGPT, or Cursor. **Temporal** orchestrates Parse → Generate → Publish. (github.com/Forge41/genalphacli)

TPS – Embedded iPaaS Platform | *Python, FastAPI, SQLAlchemy, PostgreSQL, MultiFernet*

2025 – Present

- Open-source embedded iPaaS microservice powering OAuth 2.0 / API Key / mTLS auth for AI-agent tools generated by GenAlpha, so MCP servers can hold live user credentials. Provider-abstraction layer (**AuthProvider + IntegrationService**) for pluggable third-party providers. (github.com/Forge41/genalphacli/services/tps)

TECHNICAL SKILLS

Languages & Backend: Python, TypeScript, JavaScript, Java, SQL; FastAPI, Next.js, React.js, Spring Boot, Django

GenAI & LLM Frameworks: LangChain, LangGraph, LlamaIndex, CrewAI, AutoGen, FastMCP, Anthropic SDK, OpenAI SDK, DSPy, Hugging Face Transformers

RAG & Retrieval: Retrieval-Augmented Generation, Agentic RAG, GraphRAG, hybrid search (BM25 + dense), re-ranking, semantic/recursive chunking, query rewriting; Vector DBs: Pinecone, Weaviate, Qdrant, pgvector, FAISS

Agentic AI & Orchestration: Multi-agent systems, **Model Context Protocol (MCP)**, **Google ADK**, tool use / function calling, ReAct, Plan-and-Execute, Temporal durable workflows, agent-to-agent (A2A) protocols

Voice AI: LiveKit, Pipecat, Deepgram (STT), ElevenLabs (TTS) – real-time conversational voice pipelines for agentic assistants

Evaluation, Safety & Infra: LangSmith, Langfuse, Ragas, LLM-as-a-judge, offline evals; Guardrails.ai, PII redaction, prompt-injection defense; **AWS Bedrock, Azure OpenAI, GCP Vertex AI**, vLLM, Docker, Kubernetes, CI/CD, Prometheus, Grafana, OpenTelemetry

Fine-Tuning & ML: LoRA, QLoRA, PEFT, RLHF, DPO; PyTorch, TensorFlow, scikit-learn, NLP, Computer Vision, MLOps / LLMOps

EDUCATION, CERTIFICATIONS & ADDITIONAL

Bachelor of Engineering – Computer Science & Engineering (CSE), 2020 – 2024 – Shivamogga, Karnataka. **Open Source**: Founded **Forge41** – github.com/NandishNaik01.

Anthropic Academy Certifications (Jul 2026): Introduction to Model Context Protocol · Introduction to Agent Skills · Claude Code in Action · Building with the Claude API

System Design: Sole architect of Harvey.ai's agentic workflow automation engine – HLD, LLD, production prototype in **1 week**.